

Bucky's Smart Pub

BadgerHacks 2021, Jack Wolf

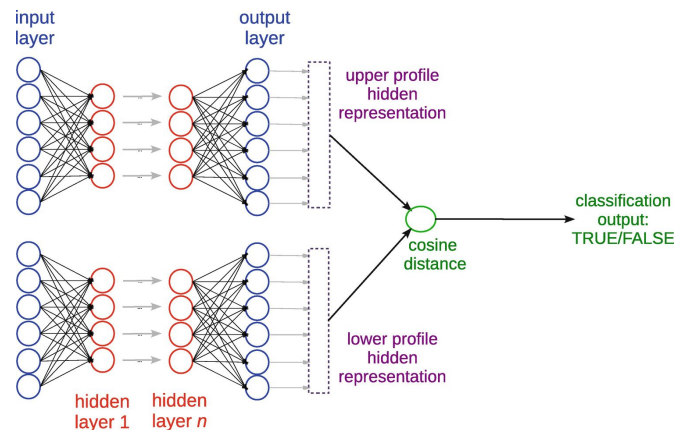
Overview

- Use ML to efficiently learn a user's preferences
- Ranking and recommendation systems
- Two parts:
 - Siamese neural network for learning pairwise ranking
 - Active learning algorithm for selecting training points
- Bucky's Smart Pub: Bucky the Bartender's goal is to learn patrons' beer preferences quickly and accurately without getting them too drunk




Siamese Neural Network

- Siamese networks use the same weights for two different inputs
- Training input:
 - Pair of points being ranked: (point_i, point_j)
 - Pairwise ranking: {+1, -1}
- Model
 - Two layers, X^\wedge and L^\wedge , of shapes (1,D) and (D,L)
 - Output of model, $L^\wedge @ X^\wedge$, is learned preference, LP
 - $\text{score}(p) = \text{norm}(P - \text{LP})$
 - $\text{Pred} = +1$ if $\text{score}(\text{point}_i) > \text{score}(\text{point}_j)$, else -1




Active Learning

- A dataset of N points has $N!(N-2)!$ Pairwise combinations
- Easy to learn preference by viewing all rankings, but not always possible
 - User cannot watch every movie on Netflix or drink every beer at the bar
- Some points carry more information than others
- Approach
 - Select initial set of training points, T
 - Put the rest of the points in a set E
 - While there are points in E :
 - $S = [\text{score}(e) \text{ for } e \text{ in } E]$
 - Selection = $E[\text{argmin}(s)]$
 - Pop selection from E to T
 - Train model(T)



Bucky's Smart Pub application

- API and website to allow users to interact with model
 - Watch progress
 - Get live recommendations
 - Give model feedback
- Add users to the pub and watch Bucky learn their preferences
- API calls ran in background so users do not have to wait



Performance

- Performance is twitchy
- Need to fix scoring heuristic

